

UPWORK PORTFOLIO CASE STUDY

Mindgard AI Security Labs

Automated Red-Teaming & Vulnerability Assessment for LLMs and AI Models

Mindgard is a London-based cybersecurity startup enabling enterprises to defend AI and LLM deployments against adversarial attacks. As Solo Designer over a 4-month sprint, I collaborated closely with a team of six (PM, 2 SWEs, AI Researcher, and Security Engineer) to translate complex neural threat telemetry into an actionable, enterprise-grade risk overview dashboard.

ROLE & SCOPE

Solo Product Designer

CLIENT / ORGANIZATION

Mindgard AI Security Labs (London, UK)

ENGAGEMENT & TIMELINE

Cross-Functional (PM, 2 SWEs, AI Researcher, SecEng)

KEY DELIVERABLES

Enterprise SaaS Dashboard, Threat Taxonomy, 1-Click Remediation

25

MITRE ATLAS Vectors

938


Threats Discovered

10

Security Experts Researched

Beta

Enterprise Beta Status



■ PROBLEM FRAMING

The Challenge & Friction Points

Enterprises deploying LLMs and machine learning models face novel adversarial attack vectors: prompt injections, jailbreaks, training data leakage, model theft, and evasion attacks. Existing cybersecurity dashboards dumped raw, probabilistic log data, leaving security teams paralyzed without actionable remediation paths or standardized severity metrics.

◆ Target User Segments & Stakeholders

Offensive Security Engineers

Conducting active penetration tests on AI models; required standardized taxonomy and repeatable scan orchestration.

AI & ML Engineers

Deploying models into production; needed clear remediation code snippets and model health benchmarks.

Enterprise CISOs & SecOps

Evaluating enterprise risk posture; needed high-level 6-month trend telemetry and executive compliance views.

■ STRATEGIC FOUNDATION

Core Design Principles

01 Simplified Risk Overview

Present executive and engineering threat postures across all models without drowning operators in raw unmanageable log dumps.

02 Historical Trend Telemetry

Track vulnerability trajectories and risk score fluctuations over 6-month horizons with comparative multi-model filters.

03 Actionable Remediation

Every vulnerability connects directly to prioritized remediation guidelines (model obfuscation, boundary hardening) with 1-click triggers.

04 MITRE ATLAS™ Standardization

Anchor threat categories directly to the industry-standard MITRE ATLAS framework with semantic color-coded risk tiers.

STRATEGIC INSIGHT


"In enterprise cybersecurity, visibility without immediate remediation breeds operational fatigue. Moving from passive log reporting to 1-click remediation transforms a monitoring tool into a proactive defense system."

■ SYSTEM BLUEPRINTS

6-Month Risk Trend & Telemetry Analytics



Multi-series risk trajectory tracking model vulnerabilities over time with customizable date ranges and model selection.



Full enterprise dashboard view showing threat categories, model inventory, and immediate mitigation actions.

◆ MITRE ATLAS Framework

Standardized 25 attack categories (Reconnaissance, Initial Access, ML Model Access, Exfiltration) into intuitive semantic donut widgets.

◆ Multi-Model Benchmarking


Engineered comparative filters allowing teams to compare Mistral, GPT-3.5, and internal proprietary models on a single graph.

◆ Frictionless Model Ingestion

Designed intuitive modal workflows for uploading local weights (.h5, .pt, .onnx) or securely connecting HuggingFace repositories.


■ INTERFACE HIGHLIGHTS

Core Product Experience & Touchpoints



● MITRE ATLAS™ Risk Donut Widget

Interactive donut chart visualizing 25 high-impact attack vectors with hover-over risk severity breakdown, active filtering, and month-over-month shifts.



● AI Model Inventory & Management

Sortable enterprise inventory tracking tested models with date last scanned, color-coded threat badges, risk gauges, attack counts, and 1-click rescan triggers.

■ SCALABLE FOUNDATIONS

Tokens, States & Parity



Tailored for technical operators and decision-makers: disciplined semantic risk tokens (Green for Low Threat 80–90%+, Amber for Medium 60%, Red for Critical 16–20%), paired with high-contrast charts, collapsible navigation, and frictionless modal dialogs.

◆ System Specifications

- ✓ Semantic risk token architecture (Low, Medium, Critical)
- ✓ Dense data tables with inline scan triggers and health badges
- ✓ Standardized modal workflows for model file and repo ingestion
- ✓ High-contrast charts optimized for enterprise dark environments

THE DESIGN ENGINEER ADVANTAGE

Bridging the gap between Figma and production front-end code (React, Next.js, TypeScript). Zero guesswork during developer handoff, pixel-perfect responsive execution, and bulletproof component architectures.



■ BUSINESS & USER IMPACT

Shipped Outcomes & Measured Results

25

MITRE ATLAS VECTORS

938

THREATS DISCOVERED

10

SECURITY EXPERTS RESEARCHED

Beta

ENTERPRISE BETA STATUS

◆ Validated Project Milestone

Delivered an end-to-end B2B AI security dashboard currently deployed in enterprise Beta testing. Integrated 25 MITRE ATLAS categories, multi-model inventory, 6-month trend telemetry, and 1-click remediation guidelines.

◆ Senior Product Reflections

"In-depth interviews with 10 offensive security engineers revealed that operators do not want more alert volume—they want prioritized remediation. Close engineering collaboration allowed us to de-scope secondary features and ship an enterprise beta that solved the core problem."

JA

Jimmy Arikawe

Senior Product Designer & Design Engineer · Ex-Pentagram

AVAILABLE ON UPWORK

◆ End-to-End Product Design (0 to 1)

◆ Design Systems & Component Libraries

◆ Complex SaaS & Enterprise Dashboards

◆ Design Engineering (React / Next.js / TS)

🌐 jimmyarikawe.com

✉️ hi@jimmyarikawe.com

📄 linkedin.com/in/jimmyarikawe

⚡ Upwork Direct Contract